

CNN Plays on Stanford Cars Dataset

TensorFlow 2.0



Introduction

Play with CNNs on a more serious classification dataset.

Data

A variant of the original Stanford Cars Dataset which was preprocessed to be a classification dataset. See details in <https://www.kaggle.com/sungtheillest/vehicledetected-stanford-cars-data-classes-folder>. Data was downloaded to x99-server at `/home/yoav/dev/other/tf2.0/data/stanford-car-dataset-by-classes-folder`.

~8000 training samples.


Training Exploration - VGG16

Doing a transfer learning with frozen weights (data is small and thus this seems like a reasonable thing to do). Transforming the input images through the VGG16 (frozen) backbone for speed. Transformed features are saved in `dev/other/tf2.0/data/stanford-car-dataset-by-classes-folder/vgg16_out`.

- Training with the following simple head:

```
x = Dense(512, activation='relu')(x)
x = Dropout(0.1)(x)
x = Dense(256, activation='relu')(x)
x = Dropout(0.1)(x)
x = Dense(256, activation='relu')(x)
x = Dropout(0.1)(x)
x = Dense(K, activation='softmax')(x)
```

Training accuracy is ~0.9 (goes to ~1.0 with no regularization); test accuracy ~0.4




There is a large difference between train and test accuracy. Since train accuracy is good (and in fact, with no regularization, the network overfits perfectly to the train

set), I tried to increase regularization: Dropout with 0.2 is reducing both accuracies; Dropout with 0.4 is not converging at all


- The first dense layer is huge with more than 10M parameters. Trying to use 1x1 convolutions to reduce the features' depth dimension (the spatial dimensions are 7x7, which cannot be significantly reduced). No improvement.....

```
i = ptm.input
x = ptm.output
x = Conv2D(64, (1, 1), activation='relu')(x)
x = Flatten()(x)
x = Dense(512, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(256, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(256, activation='relu')(x)
x = Dropout(0.2)(x)
x = Dense(K, activation='softmax')(x)
```

- With dropout of 0.1, and with decaying learning rate (), got slightly better results: train accuracy: 0.9877 , val_accuracy: 0.4622
- Trained with data augmentation (not using the transformed features), is much much slower, but is improving the test accuracy dramatically. See plot below after 70 epochs:



After more training (total of 300 epochs) this is what we get:




Training Exploration - Inception V3

Transformed features are saved in `dev/other/tf2.0/data/stanford-car-dataset-by-classes-folder/inception_v3_out`.

Inception V3 with the head above (with 1x1 conv of 64 channels) got ~ 30% test accuracy

With data augmentation the result is even worse -

Model InceptionV3



Training Exploration - ResNet50 V2

Transformed features are saved in dev/other/tf2.0/data/stanford-car-dataset-by-classes-folder/resnetv2_out.

ResNet50 V2 with the head above (with 1x1 conv of 64 channels) got ~ 35% test accuracy

No improvement when training with data augmentation

On the larger train2 training set (50 epochs): accuracy: 0.781; val_accuracy: 0.369

Next Steps

Summary:

VGG16 based network with data augmentation produces the best results - ~0.6 test accuracy. Checkpoints for best model are in
~/dev/other/tf2.0/training_checkpoints/stanford-car-dataset/vgg16/1/

- Increase train set:
 - Increase train set by 50% on the expense of the test set.
 - New folders: train2, test2 with sizes 11400, 4700 respectively (before they were approximately equal)